

ECG Arrhythmia Classification

Technical Project Report — Production-Ready Machine Learning Pipeline for Physiological Signal Analysis

PYTHON • SCIKIT-LEARN • SCIPY

STATUS: PRODUCTION-READY

Author	Joti Narinder 3rd-year Bachelor Student in Biomedical Civil Engineering University of Liège (ULiège)	Date	July 2026
Technical Stack	Python, wfdb, SciPy, Scikit-Learn, Matplotlib, Seaborn	Dataset source	MIT-BIH Arrhythmia Database (PhysioNet)
Project Domain	Biomedical Engineering, Automated Cardiac Rhythm Diagnostics, Feature Engineering.		

1. Executive Summary & Context

In the rapidly evolving MedTech landscape, automated processing of Electrocardiogram (ECG) data is essential for scalable clinical decision support systems. This project presents a robust, reproducible end-to-end Python pipeline designed to ingest raw, noisy physiological signals, isolate individual heartbeats, and systematically classify them into healthy (Normal Sinus Rhythm) or pathological states (Arrhythmias).

Engineered to meet production-level coding and documentation standards, this repository serves as a concrete proof of concept demonstrating core competencies in physiological signal processing, custom statistical feature extraction, and traditional machine learning classification metrics.

2. Dataset Architecture & Ingestion

The pipeline relies on the gold-standard **MIT-BIH Arrhythmia Database** hosted on **PhysioNet**. This database contains annotated continuous ambulatory ECG recordings validated by expert cardiologists.

- **Dataset Scale:** 21,561 fully segmented individual heartbeats extracted across 10 extensive clinical records.
- **Signal Specification:** Continuous two-channel ambulatory ECG signals sampled at 360 Hz.
- **Target Classes:** Binary classification mapping normal beats against abnormal ventricular arrhythmias (Premature Ventricular Contractions).

3. End-to-End Pipeline Engineering

The architecture follows a strict modular design pattern to ensure high maintainability and computational traceability:

Phase 1: Native Ingestion via Waveform Database (`wfdb`)

To preserve signal fidelity, raw binary data and expert native annotations (`.atr`) are ingested programmatically using the specialized `wfdb` API. This step completely eliminates errors associated with intermediate plain-text or CSV conversions.

Phase 2: Digital Signal Processing & Noise Reduction (`scipy.signal`)

Raw ECG recordings are heavily prone to high-frequency muscle artifacts and low-frequency baseline wander (caused by patient breathing). A zero-phase **Butterworth bandpass filter (0.5 - 50 Hz)** was implemented to aggressively clean the waveform, isolating the relevant QRS complex while maintaining signal phase alignment.

Phase 3: R-Peak Localization & Standardized Segmentation

Using localized amplitude peaks, the pipeline maps the exact temporal coordinates of the R-waves. A fixed-size temporal window is centered around each R-peak, isolating each individual heartbeat contraction into an independent matrix slice.

Phase 4: Statistical Feature Engineering


Rather than directly feeding high-dimensional raw signal windows into the classifier, an explicit feature engineering method was built. Eight descriptive statistical metrics are extracted per window: **Mean, Variance, Minimum, Maximum, Median, Skewness, Kurtosis**, and **Peak Amplitude**. This optimizes compute efficiency and increases model stability.

Phase 5: Classification Modeling (`scikit-learn`)

A **Random Forest Classifier** was trained to classify the resulting feature matrix. This algorithm was selected due to its innate resilience to overfitting, high performance on tabular datasets, and built-in capability to compute feature importances, enabling post-hoc explainability.

4. Results & Critical Performance Evaluation

The trained system was rigorously evaluated on a pristine hold-out testing partition containing **4,313 independent heartbeats**.



Detailed Confusion Matrix Analysis

The fine-grained test results are detailed below:

Actual \ Predicted Class	Predicted: Normal	Predicted: Arrhythmia	Class Success Rate
Actual: Normal (4,179)	4,174 (True Negatives)	5 (False Positifs)	99.88%
Actual: Arrhythmia (134)	30 (False Negatives)	104 (True Positives)	77.61%

Clinical Interpretation of False Negatives:

The model demonstrates near-flawless precision when identifying healthy cardiac rhythms (only 5 false alarms out of over 4,100 instances). However, 30 pathological beats bypassed the system (classified as normal). Morphological overlay validation indicates that these specific 30 cases have borderline structural shapes that look almost identical to healthy rhythms, misleading a standard linear statistical classifier.

Feature Importance & Model Explainability

The tree-based split calculations show that the **Minimum** and **Median** values of the heartbeat window concentrate the largest share of predictive power. Physiologically, this indicates that the deeper, lower contours following the R-peak carry the strongest statistical signature of an arrhythmia when using tabular architectures.

5. Technical Roadmaps for Next-Gen Upgrades (Phase 2)

To surpass the current limitation of 30 false negatives and achieve strict clinical deployment standards, two upgrade paths are planned:

1. **Transition to Deep Learning (1D-CNN):** Implementing a 1-Dimensional Convolutional Neural Network directly on the filtered time-series signal. This completely skips manual feature engineering. Convolutional kernels automatically learn micro-spatial patterns, capturing the complex wave changes that traditional statistical metrics miss.
2. **Class Imbalance Mitigation:** Incorporating synthetic oversampling techniques (SMOTE) or tuning the classifier to enforce `class_weight='balanced'` to penalize false negatives on the minority class more heavily.

6. Conclusion

This pipeline achieves an exceptional global accuracy of **99.2%**, successfully validating classical machine learning models for production workflows on physiological signals. The code structure, clean architecture, and thorough reporting stand as evidence of a disciplined, engineering-first approach to health-tech artificial intelligence.